

# CS4012 Topics in Functional Programming

Glenn Strong

## Parallel Programming in Haskell

- Parallelism has been talked about in functional programming circles since the earliest days
- Immutable data and side-effect free computation has a lot of promise

### A small example


```
1  fib :: Integer -> Integer
2  fib n | n < 2 = 1
3  fib n = fib (n-1) + fib (n-2)
4
5  main = print $ fib 37
```

### Compile:

```
ghc -threaded -rtsopts -eventlog ex1.hs
```

### Sample runs:

```
$ ./ex1
$ ./ex1 +RTS -N1      # 1 Core
$ ./ex1 +RTS -N4      # 4 Cores
$ ./ex1 +RTS -N4 -ls  # 4 Cores + event log
```




## Event logs

We can look at the event log using the ThreadScope tool

- demo time!
- ThreadScope can be a bit of a pain to install
- Use a prebuild binary if you can

## ThreadScope main view



## Results!

That's not so good. Let's try *telling* the run-time where some parallelism is possible.

## “Spark” parallelism

The `par` operation (from the `Control.Parallel` library) allows us to signal sites of potential parallelism.

```
par :: a -> b -> b
```

Semantically `par x y` is equivalent to just `y`, but the runtime is allowed to use it as a hint.

Next try

```
1 import Control.Parallel
2
3 fib :: Integer -> Integer
4 fib n | n < 2 = 1
5 fib n = par nf ( fib (n-1) + nf )
6         where nf = fib (n-2)
7
8 main = print $ fib 37
```

## Timing

It actually got slower!

## Threadscope

It looks OK at first...

## Threadscope

Zooming in shows lots of threads stalling...

## What's happening?

```
5 fib n = par nf ( fib (n-1) + nf )
6         where nf = fib (n-2)
```




Figure 1: Runtime vs HEC's




Figure 2: ThreadScope(example 1, 4 cores)



Figure 3: ThreadScope(example 1, 4 cores)

There is a new task for each `nf`; one starts and then blocks on the other right away!

The order that `(+)` evaluates it's arguments is at the heart of this.

## Taking over a bit

So if we just swap around the arguments would that be enough?

```
5 fib n = par nf (nf + fib (n-2) )
6       where nf = fib (n-1)
```

That's crazy!

A better way...

We shouldn't have to care about how `(+)` treats its' arguments.

Introducing:

```
pseq :: a -> b -> b
```

## Using pseq

Evaluate `x` *before* `y`, returning `y`.

```
1 fib n | n < 2 = 1
2 fib n = par nf1 (pseq nf2 (nf1 + nf2))
3       where nf1 = fib (n-1)
4             nf2 = fib (n-2)
```

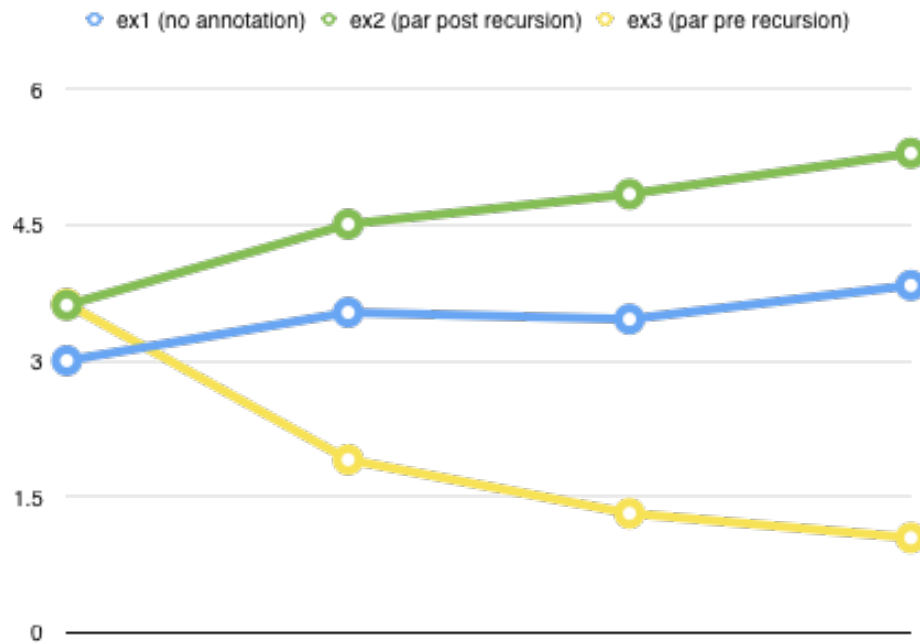


Figure 4: Runtime vs HEC's

## Further tweaking

Spark overhead can dominate after a while. Limit new threads to allow a more even distribution of work?

```

1 import Control.Parallel
2
3 -- Sequential version of fib, when we want to avoid parallelism
4 sfib :: Integer -> Integer
5 sfib n | n < 2 = 1
6 sfib n = sfib (n-1) + sfib (n-2)
7
8 fib :: Integer -> Integer -> Integer
9 fib 0 n = sfib n
10 fib _ n | n < 2 = 1
11 fib d n = par nf1 (pseq nf2 (nf1 + nf2))
12         where nf1 = fib (d-1) (n-1)
13               nf2 = fib (d-1) (n-2)
14
15 main = print $ fib 3 37

```

## Going further

- This is all still a bit tricky to use
- Lots of things to think about
  - Evaluated vs. Unevaluated computation
  - relative costs and sizes of computation
  - Sharing
- To explore how Haskell addresses these we need to take a digression...